HepaBot

Group members:
Hasnain Sohail 24280029

Scope and Deliverables

🎓 **Project Title:**
HepaBot – An AI-powered Voice Assistant for Liver Diagnostics and Report Generation

📝 **Brief Description:**
Streamlines clinical documentation by transcribing doctor-patient conversations, performing speaker and role identification, and auto-generating structured medical reports.

📊 **System Overview:**

- **Input:** Audio from real-time patient consultation or uploaded recordings
- **Processing:** Whisper for transcription → Pyannote for diarization → ClinicalBERT for role classification
- **Output:** Structured, labeled medical report (D: Doctor / P: Patient) in .txt / .csv formats

✨ **What Makes It Unique:**

- Context-aware retrieval from fragmented medical speech
- Structured doctor-patient dialogue generation
- Role classification using fine-tuned ClinicalBERT
- Granular section-level search (symptoms, diagnoses, treatments)
- Auto-generated clinical summaries from fragmented records
- AI enhanced queries

🎯 **Intended Audience:**

- Clinical QA teams
- Medical researchers
- Hospital digitization staff

Project Progress

📦 **Deliverables & Completion:**

- ✔ Transcription (Whisper) – 100%
- 🔄 Speaker Diarization – 50%
- 🔄 PDF-to-Structured Pipeline – 70%
- 🔍 Vector Search & Retrieval – 40%
- 📈 Analytics Dashboard – 70%
- 📝 Auto-Report Generation – 80%

📌 **Current Status:**

- ✔ Working: Audio pipeline, basic semantic search, role classification, report generation
- ⚠ Blockers: Custom model inference on Hugging Face, UML-based entity extraction

🛠 **Tech Stack:**

- **Frameworks:** Streamlit, LangChain, ChromaDB
- **Models:** Whisper, pyannote.audio v3.x, ClinicalBERT (fine-tuned), LLaMA-3 (via Ollama)
- **Languages & Libraries:** PyTorch, HuggingFace Transformers, Pandas, Scikit-learn

📏 **Evaluation Metrics:**

- Transcription WER (planned)
- Accuracy for role classification (TBD)
- Latency for real-time pipeline

Live Demo Highlights

💻 **What You'll Showcase Live:**

1. Audio-to-Transcript flow using Whisper
2. Speaker Diarization results
3. Role Identification via ClinicalBERT
4. Downloadable D/P Structured Report
5. Summary Generation
6. Structured Json creation and storing in Vector DB
7. Patient Analytics
8. Semantic Search (Partially)
9. Enhanced Search capabilities with AI

📸 **Key Outputs / UI Elements:**

- 🎙️ Transcription text area
- 🎭 Speaker-labeled segments
- 🧑‍⚕️ Role-tagged final dialogue (D: / P:)
- ⬇️ CSV and TXT report download

📯 **Technical Wins:**

- ClinicalBERT role detection on noisy transcripts
- Pyannote.audio v3.x migration (modern pipeline adaptation)
- Whisper performance optimization on long consultations
- Hybrid retrieval using both vector search and structured metadata
- Fine-tunning Llama2 7b on medical dataset using PEFT

What’s Next:

- Improve diarization speaker-label accuracy
- Refine prompt tuning for LLaMA-3 to enhance semantic search
- Deploy on Streamlit Cloud / Hugging Face Spaces

⚙ **Solving Blockers / Boosting Performance:**

- Enable gated model access with HuggingFace tokens
- Entity extraction via MedSpaCy or UML graph embeddings
- Caching for inference speed-up and latency reduction

Hosting Plan:

- Target: Streamlit Cloud for lightweight deployment
- Backup: Local Dockerized version for hospitals without public cloud access

Dependencies / Support Needed:

- HuggingFace token access to pyannote/speaker-diarization
- Guidance on advanced speaker-role assignment logic (multi-speaker cases)

SCREENSHOTS

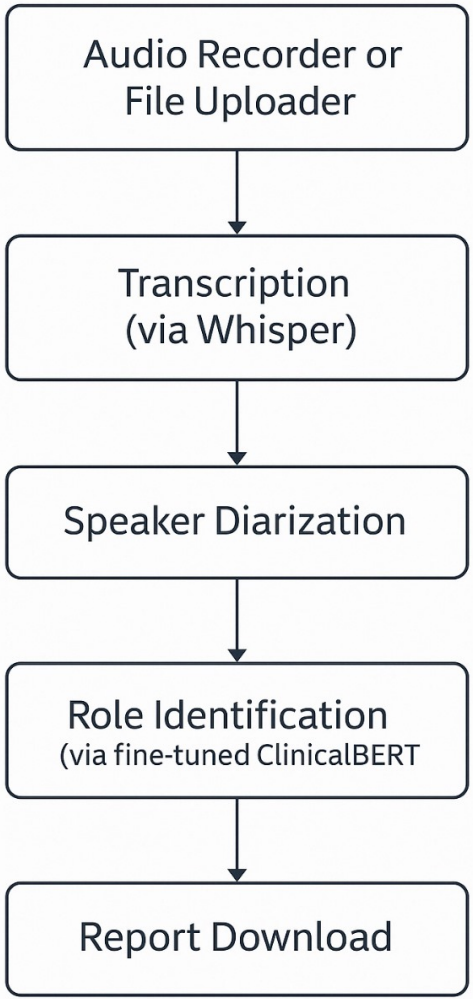
Audio Recorder or
File Uploader
Transcription
(via Whisper)
Speaker Diarization
Role Identification
(via fine-tuned ClinicalBERT
Report Download

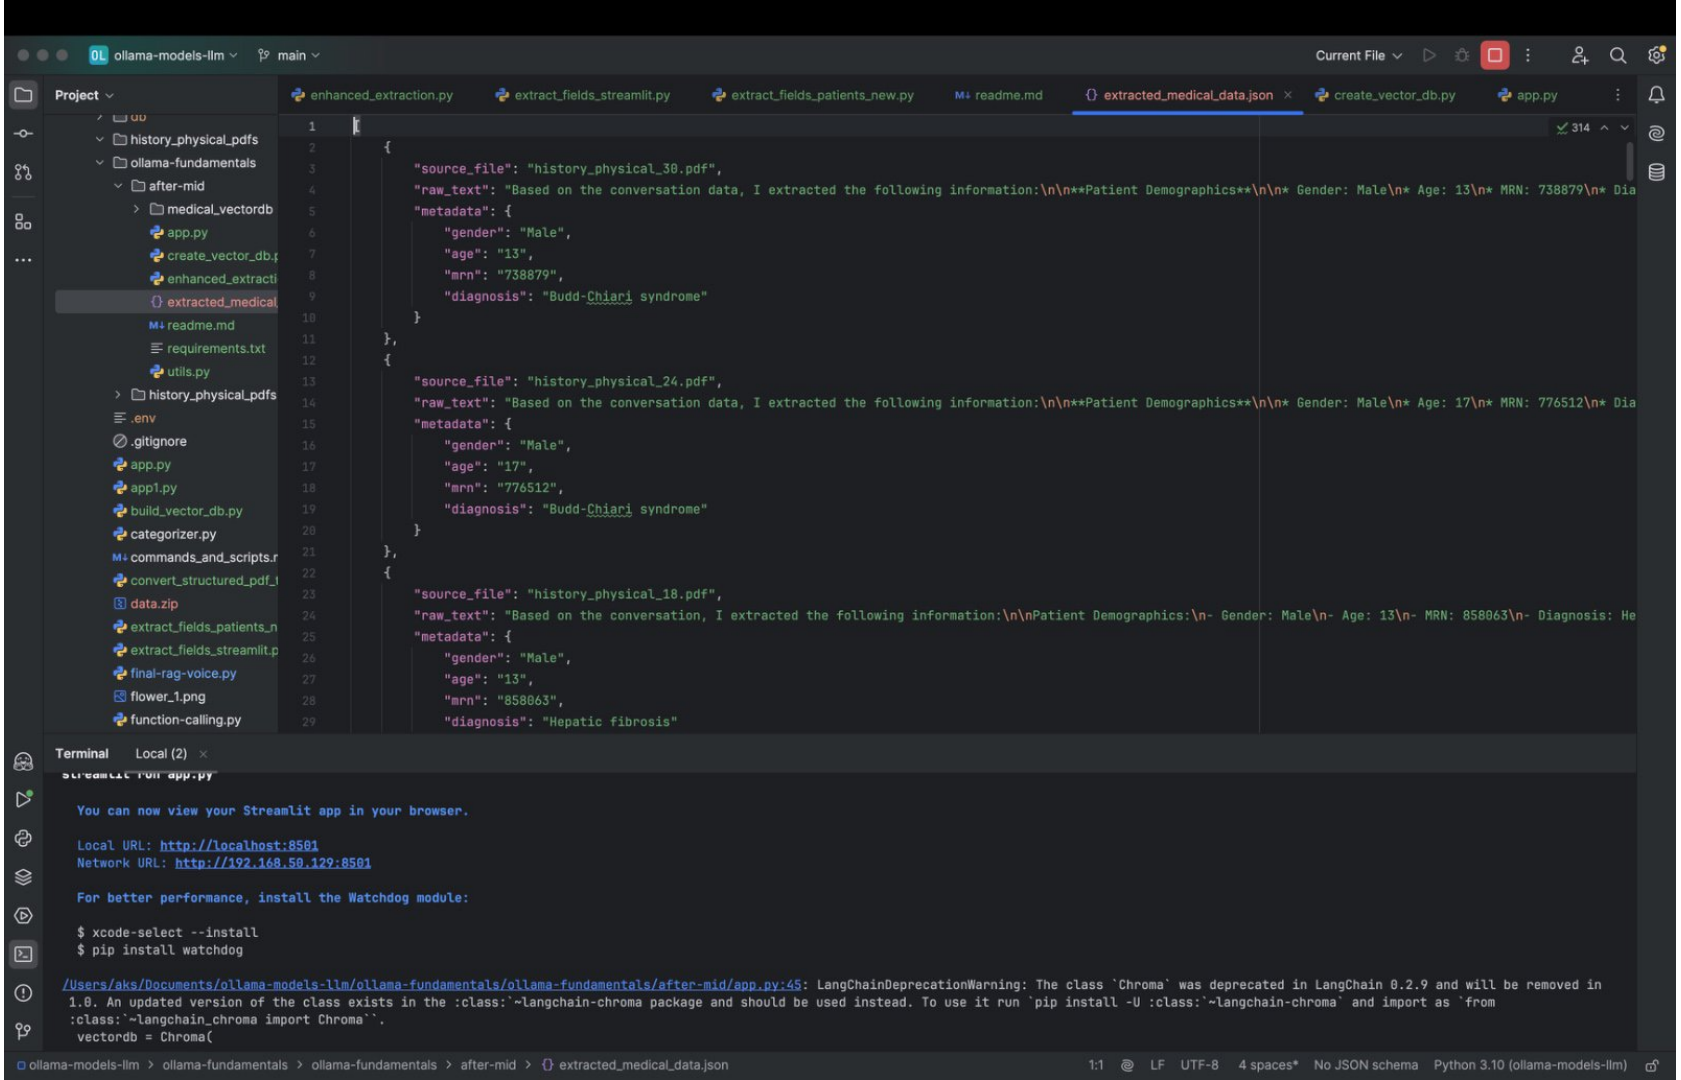
ollama-models-llm
main
Current File
Project
enhanced_extraction.py
extract_fields_streamlit.py
extract_fields_patients_new.py
readme.md
extracted_medical_data.json
create_vector_db.py
app.py
history_physical_pdfs
ollama-fundamentals
after-mid
medical_vectordb
app.py
create_vector_db.p
enhanced_extracti
extracted_medical
readme.md
requirements.txt
utils.py
history_physical_pdfs
.env
.gitignore
app.py
app1.py
build_vector_db.py
categorizer.py
commands_and_scripts.r
convert_structured_pdf_t
data.zip
extract_fields_patients_n
extract_fields_streamlit.p
final-rag-voice.py
flower_1.png
function-calling.py
[
{
"source_file": "history_physical_30.pdf",
"raw_text": "Based on the conversation data, I extracted the following information:\n\n**Patient Demographics**\n\n* Gender: Male\n* Age: 13\n* MRN: 738879\n* Dia
"metadata": {
"gender": "Male",
"age": "13",
"mrn": "738879",
"diagnosis": "Budd-Chiari syndrome"
}
},
{
"source_file": "history_physical_24.pdf",
"raw_text": "Based on the conversation data, I extracted the following information:\n\n**Patient Demographics**\n\n* Gender: Male\n* Age: 17\n* MRN: 776512\n* Dia
"metadata": {
"gender": "Male",
"age": "17",
"mrn": "776512",
"diagnosis": "Budd-Chiari syndrome"
}
},
{
"source_file": "history_physical_18.pdf",
"raw_text": "Based on the conversation, I extracted the following information:\n\nPatient Demographics:\n- Gender: Male\n- Age: 13\n- MRN: 858063\n- Diagnosis: He
"metadata": {
"gender": "Male",
"age": "13",
"mrn": "858063",
"diagnosis": "Hepatic fibrosis"
Terminal
Local (2)
You can now view your Streamlit app in your browser.
Local URL: http://localhost:8501
Network URL: http://192.168.50.129:8501
For better performance, install the Watchdog module:
$ xcode-select --install
$ pip install watchdog
/Users/aks/Documents/ollama-models-llm/ollama-fundamentals/ollama-fundamentals/after-mid/app.py:45: LangChainDeprecationWarning: The class `Chroma` was deprecated in LangChain 0.2.9 and will be removed in
1.0. An updated version of the class exists in the :class:`~langchain-chroma package and should be used instead. To use it run `pip install -U :class:`~langchain-chroma` and import as `from
:class:`~langchain_chroma import Chroma``.
vectordb = Chroma(
ollama-models-llm > ollama-fundamentals > ollama-fundamentals > after-mid > extracted_medical_data.json
1:1
LF
UTF-8
4 spaces*
No JSON schema
Python 3.10 (ollama-models-llm)

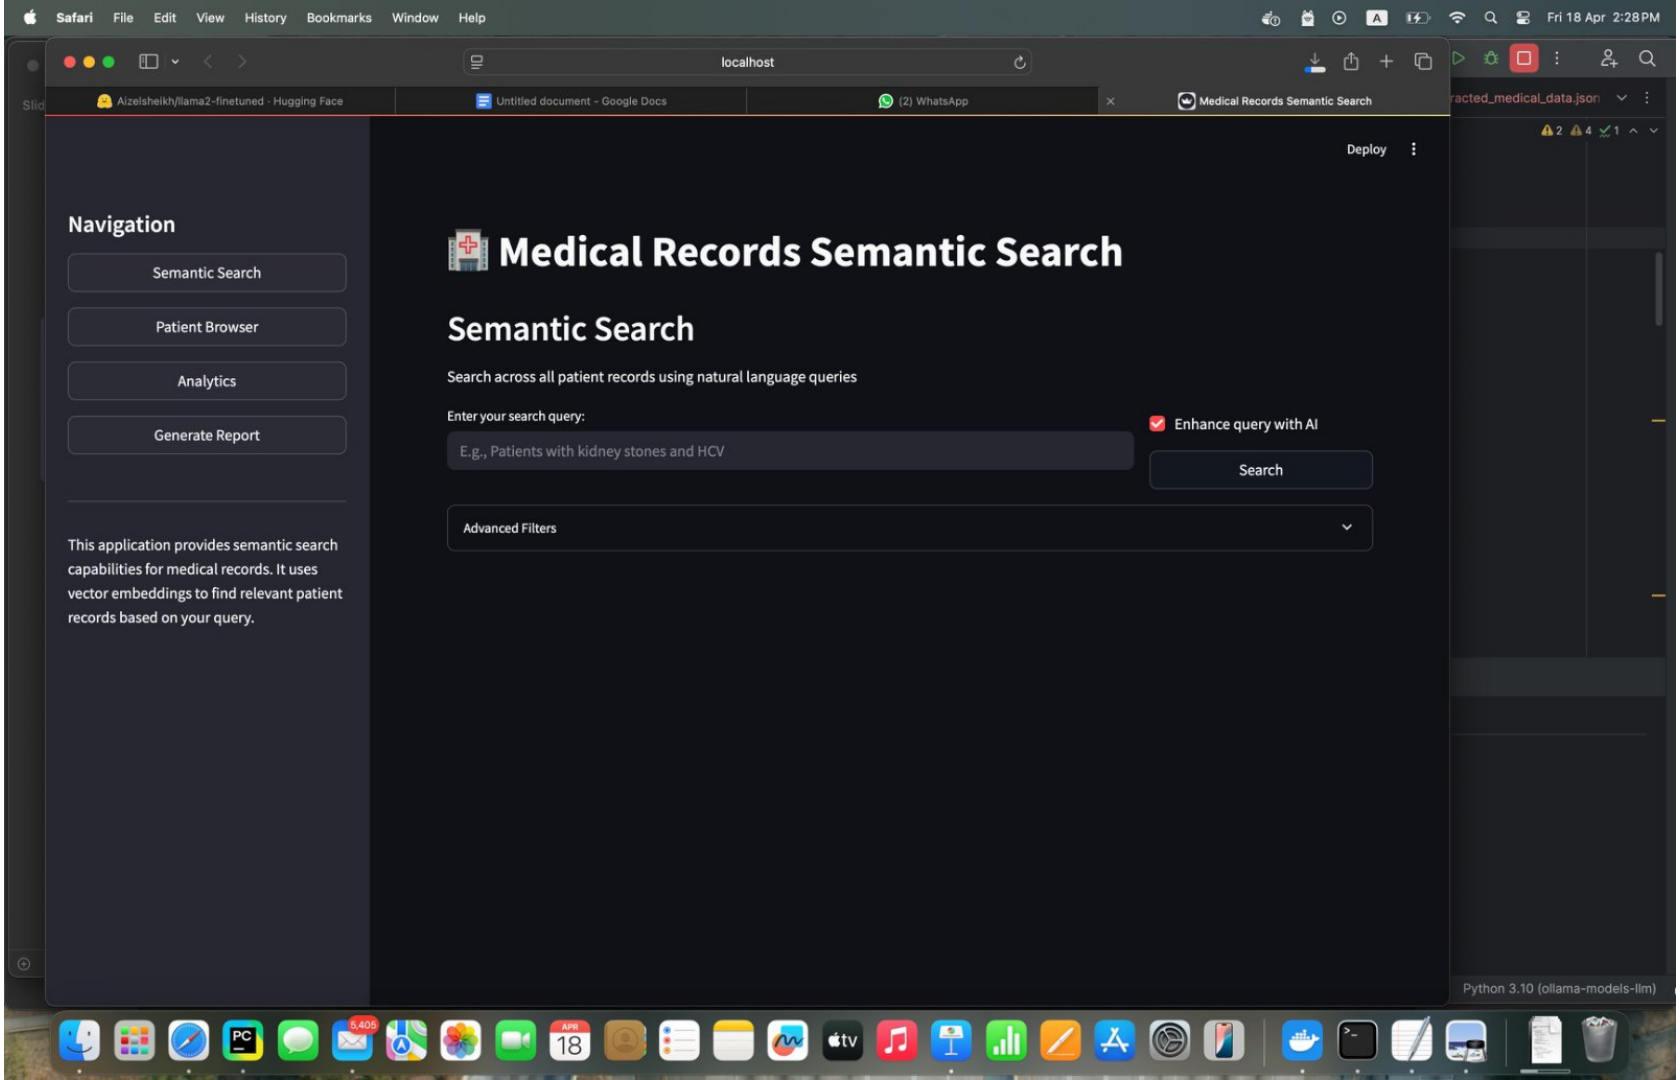
Safari File Edit View History Bookmarks Window Help
Fri 18 Apr 2:28 PM
localhost
Aizelsheikh/llama2-finetuned · Hugging Face
Untitled document - Google Docs
(2) WhatsApp
Medical Records Semantic Search
Deploy
Navigation
Semantic Search
Patient Browser
Analytics
Generate Report
This application provides semantic search capabilities for medical records. It uses vector embeddings to find relevant patient records based on your query.
Medical Records Semantic Search
Semantic Search
Search across all patient records using natural language queries
Enter your search query:
E.g., Patients with kidney stones and HCV
Enhance query with AI
Search
Advanced Filters
Python 3.10 (ollama-models-llm)

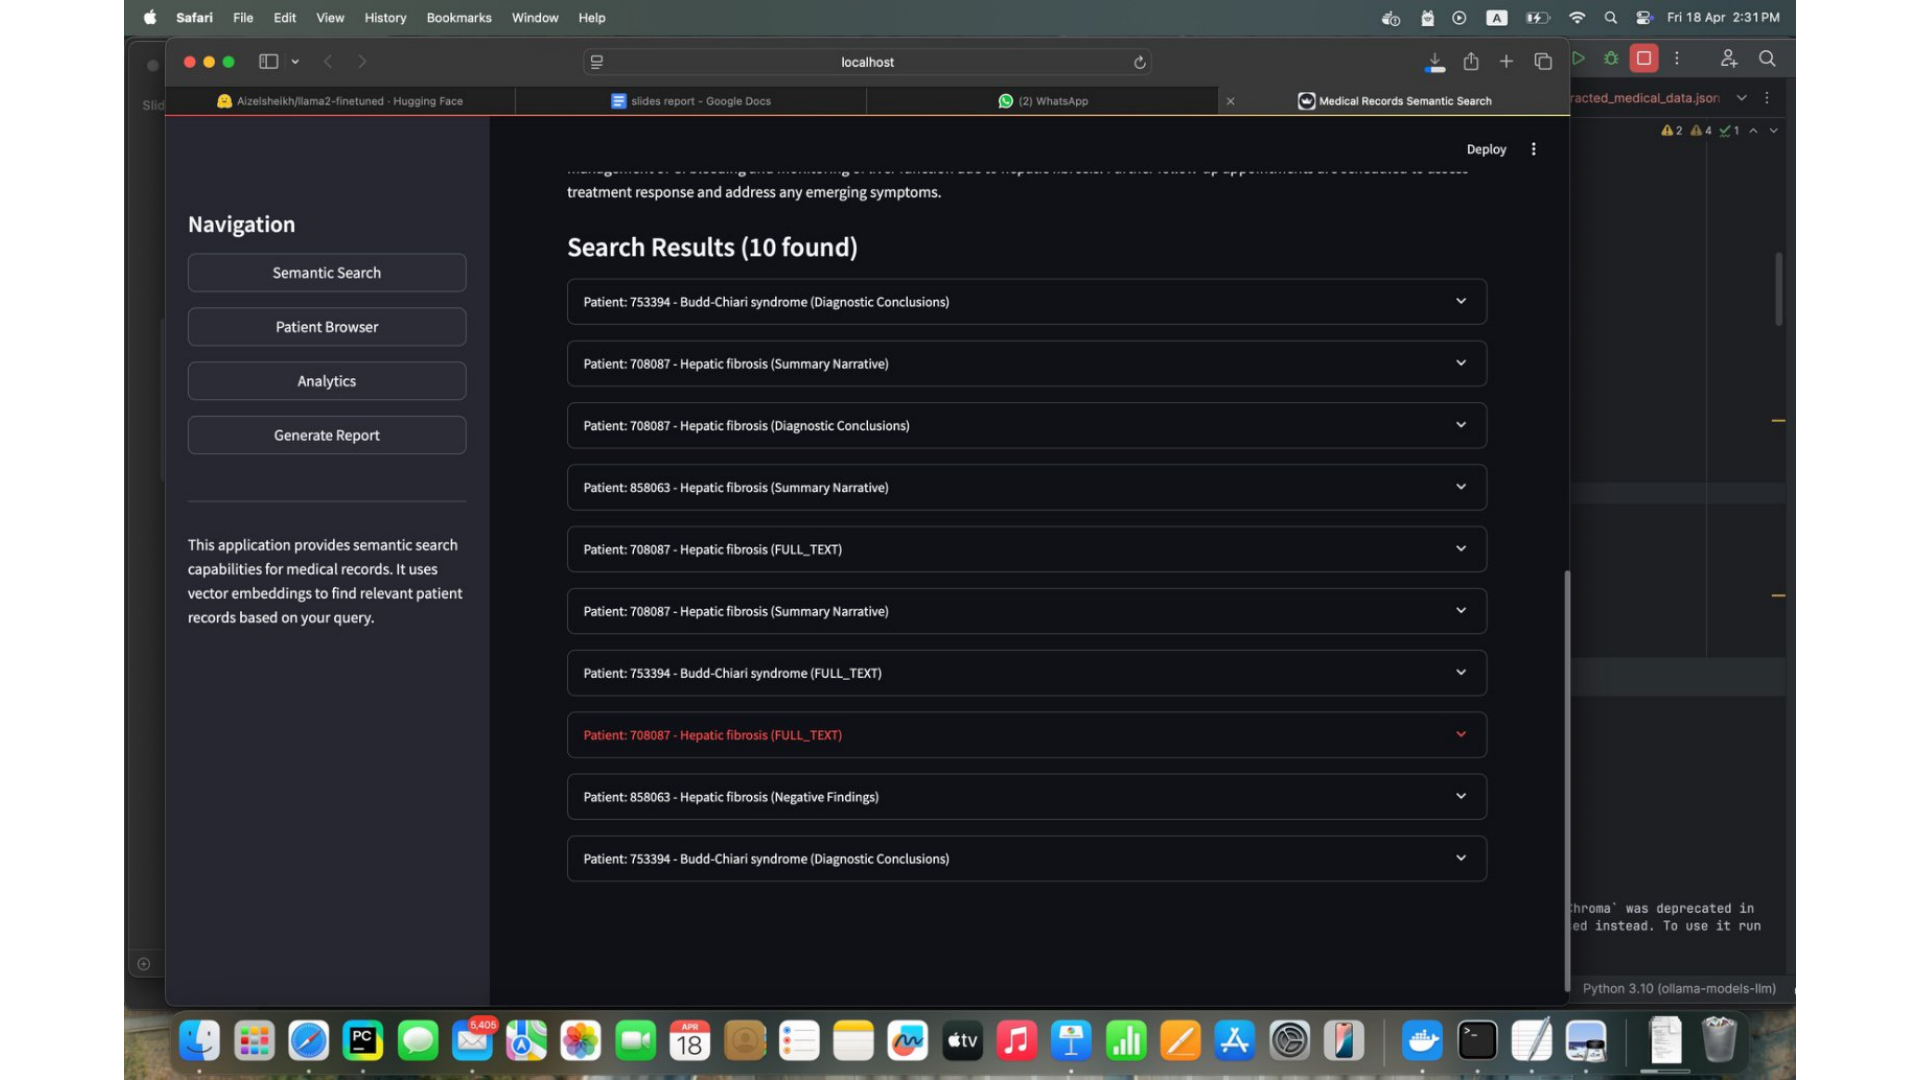
Safari File Edit View History Bookmarks Window Help
Fri 18 Apr 2:31 PM
localhost
Aizelsheikh/llama2-finetuned · Hugging Face
slides report - Google Docs
(2) WhatsApp
Medical Records Semantic Search
Deploy
Navigation
Semantic Search
Patient Browser
Analytics
Generate Report
This application provides semantic search capabilities for medical records. It uses vector embeddings to find relevant patient records based on your query.
treatment response and address any emerging symptoms.
Search Results (10 found)
Patient: 753394 - Budd-Chiari syndrome (Diagnostic Conclusions)
Patient: 708087 - Hepatic fibrosis (Summary Narrative)
Patient: 708087 - Hepatic fibrosis (Diagnostic Conclusions)
Patient: 858063 - Hepatic fibrosis (Summary Narrative)
Patient: 708087 - Hepatic fibrosis (FULL_TEXT)
Patient: 708087 - Hepatic fibrosis (Summary Narrative)
Patient: 753394 - Budd-Chiari syndrome (FULL_TEXT)
Patient: 708087 - Hepatic fibrosis (FULL_TEXT)
Patient: 858063 - Hepatic fibrosis (Negative Findings)
Patient: 753394 - Budd-Chiari syndrome (Diagnostic Conclusions)
Python 3.10 (ollama-models-llm)

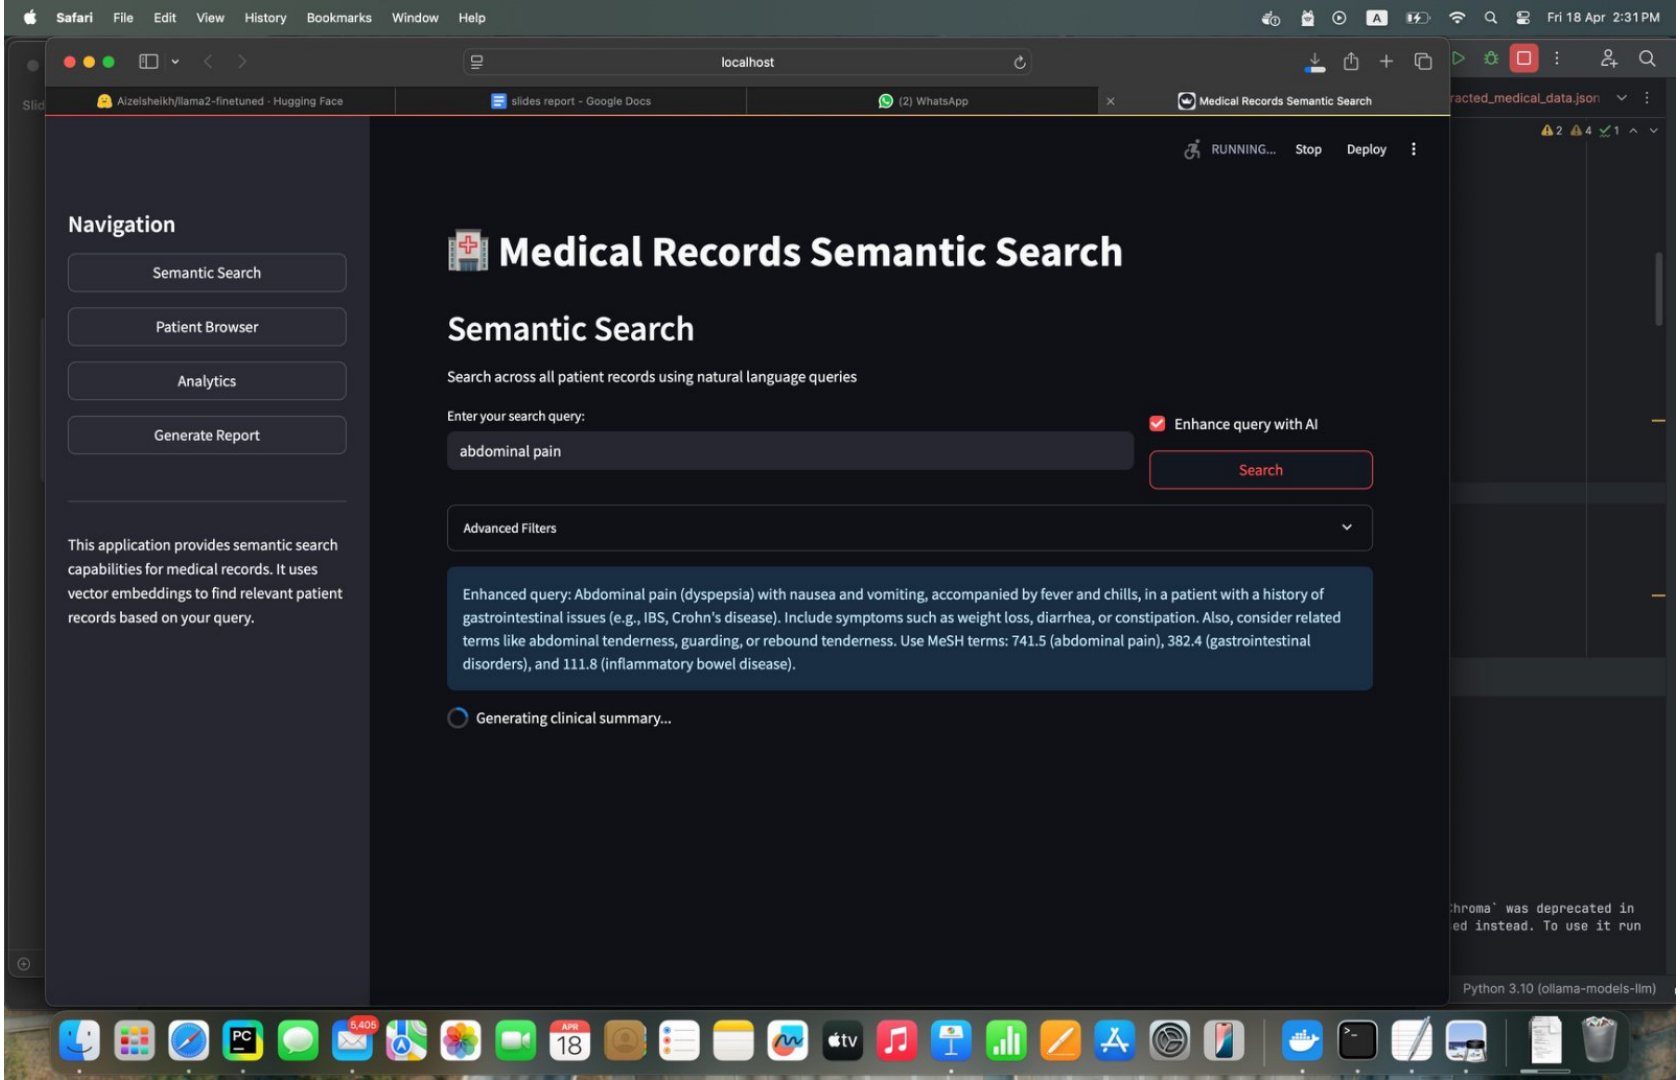
Safari File Edit View History Bookmarks Window Help
Fri 18 Apr 2:31 PM
localhost
Aizelsheikh/llama2-finetuned · Hugging Face
slides report - Google Docs
(2) WhatsApp
Medical Records Semantic Search
RUNNING... Stop Deploy
Navigation
Semantic Search
Patient Browser
Analytics
Generate Report
This application provides semantic search capabilities for medical records. It uses vector embeddings to find relevant patient records based on your query.
Medical Records Semantic Search
Semantic Search
Search across all patient records using natural language queries
Enter your search query:
abdominal pain
Enhance query with AI
Search
Advanced Filters
Enhanced query: Abdominal pain (dyspepsia) with nausea and vomiting, accompanied by fever and chills, in a patient with a history of gastrointestinal issues (e.g., IBS, Crohn's disease). Include symptoms such as weight loss, diarrhea, or constipation. Also, consider related terms like abdominal tenderness, guarding, or rebound tenderness. Use MeSH terms: 741.5 (abdominal pain), 382.4 (gastrointestinal disorders), and 111.8 (inflammatory bowel disease).
Generating clinical summary...
Python 3.10 (ollama-models-llm)

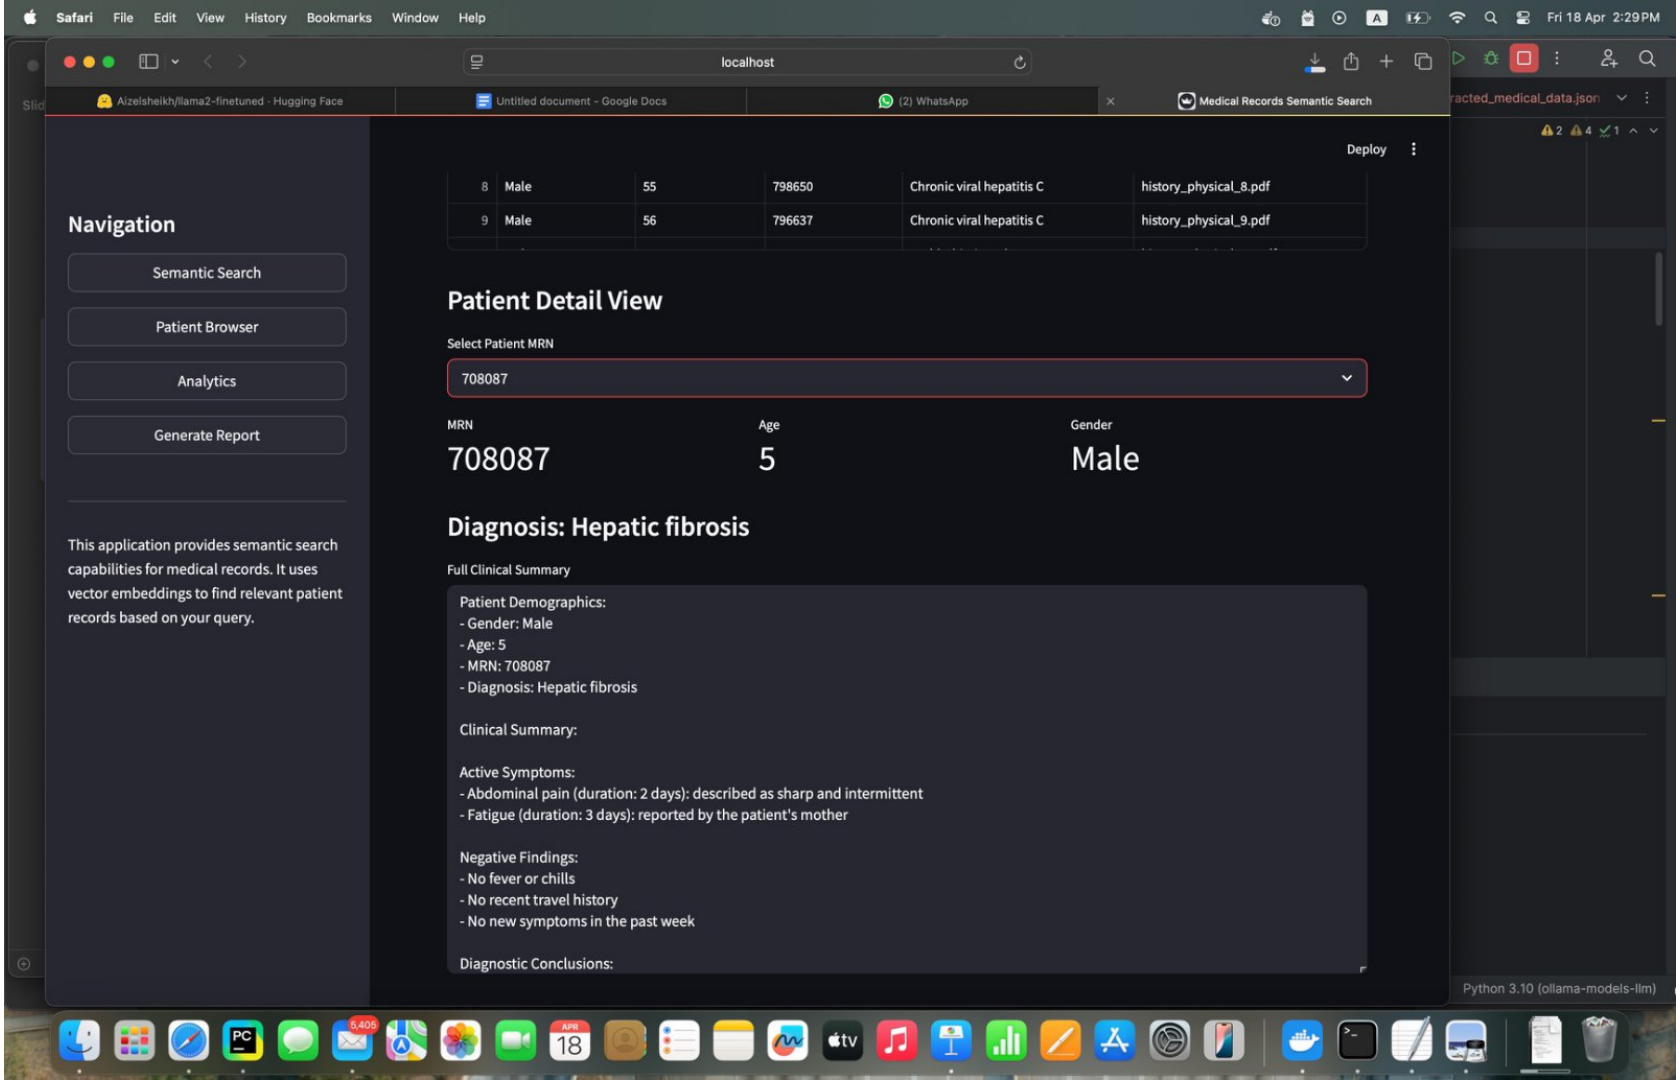
Safari File Edit View History Bookmarks Window Help
Fri 18 Apr 2:29 PM
localhost
Aizelsheikh/llama2-finetuned · Hugging Face
Untitled document - Google Docs
(2) WhatsApp
Medical Records Semantic Search
Deploy
Navigation
Semantic Search
Patient Browser
Analytics
Generate Report
This application provides semantic search capabilities for medical records. It uses vector embeddings to find relevant patient records based on your query.
8 Male 55 798650 Chronic viral hepatitis C history_physical_8.pdf
9 Male 56 796637 Chronic viral hepatitis C history_physical_9.pdf
Patient Detail View
Select Patient MRN
708087
MRN
708087
Age
5
Gender
Male
Diagnosis: Hepatic fibrosis
Full Clinical Summary
Patient Demographics:
- Gender: Male
- Age: 5
- MRN: 708087
- Diagnosis: Hepatic fibrosis
Clinical Summary:
Active Symptoms:
- Abdominal pain (duration: 2 days): described as sharp and intermittent
- Fatigue (duration: 3 days): reported by the patient's mother
Negative Findings:
- No fever or chills
- No recent travel history
- No new symptoms in the past week
Diagnostic Conclusions:
Python 3.10 (ollama-models-llm)

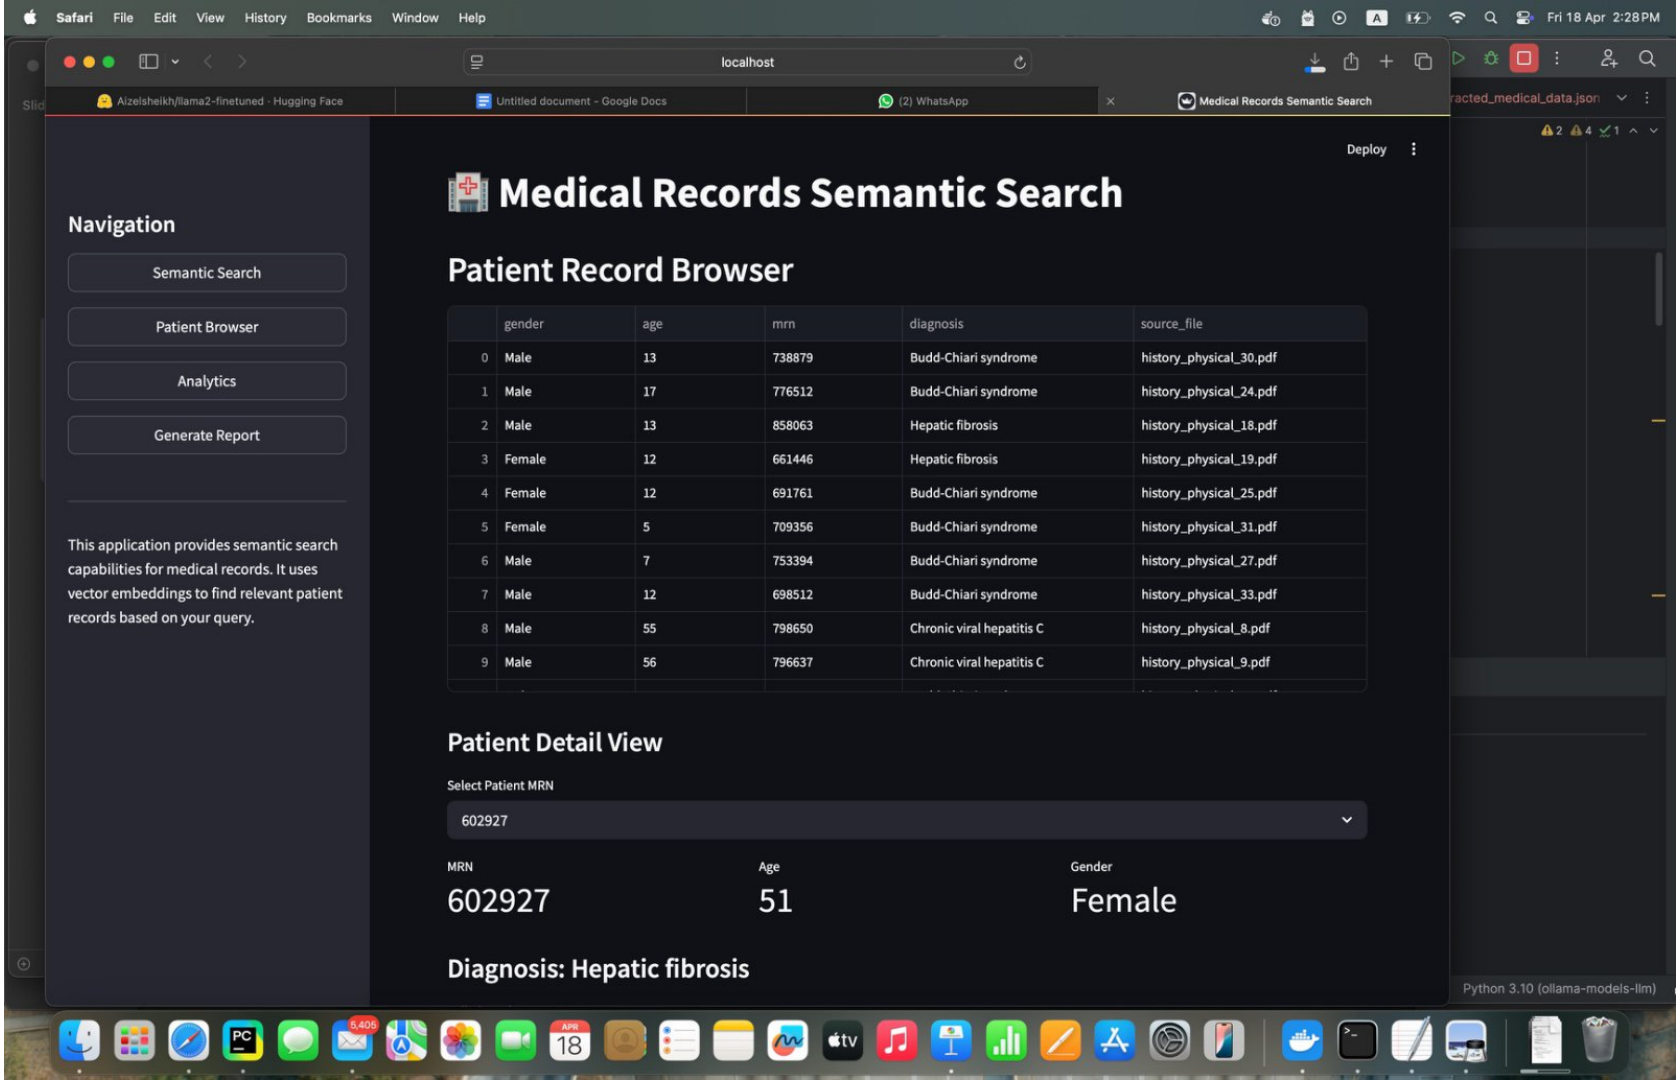
Safari File Edit View History Bookmarks Window Help
Fri 18 Apr 2:28 PM
localhost
Aizelsheikh/llama2-finetuned · Hugging Face
Untitled document - Google Docs
(2) WhatsApp
Medical Records Semantic Search
Deploy
Medical Records Semantic Search
Navigation
Semantic Search
Patient Browser
Analytics
Generate Report
This application provides semantic search capabilities for medical records. It uses vector embeddings to find relevant patient records based on your query.
Patient Record Browser
gender age mrn diagnosis source_file
0 Male 13 738879 Budd-Chiari syndrome history_physical_30.pdf
1 Male 17 776512 Budd-Chiari syndrome history_physical_24.pdf
2 Male 13 858063 Hepatic fibrosis history_physical_18.pdf
3 Female 12 661446 Hepatic fibrosis history_physical_19.pdf
4 Female 12 691761 Budd-Chiari syndrome history_physical_25.pdf
5 Female 5 709356 Budd-Chiari syndrome history_physical_31.pdf
6 Male 7 753394 Budd-Chiari syndrome history_physical_27.pdf
7 Male 12 698512 Budd-Chiari syndrome history_physical_33.pdf
8 Male 55 798650 Chronic viral hepatitis C history_physical_8.pdf
9 Male 56 796637 Chronic viral hepatitis C history_physical_9.pdf
Patient Detail View
Select Patient MRN
602927
MRN
602927
Age
51
Gender
Female
Diagnosis: Hepatic fibrosis
Python 3.10 (ollama-models-llm)

Safari File Edit View History Bookmarks Window Help

Fri 18 Apr 2:28 PM

localhost

Aizelsheikh/llama2-finetuned · Hugging Face | Untitled document - Google Docs | (2) WhatsApp | Medical Records Semantic Search

Deploy

Medical Records Semantic Search

Navigation

- Semantic Search
- Patient Browser
- Analytics
- Generate Report

This application provides semantic search capabilities for medical records. It uses vector embeddings to find relevant patient records based on your query.

Medical Records Analytics

Patient Demographics

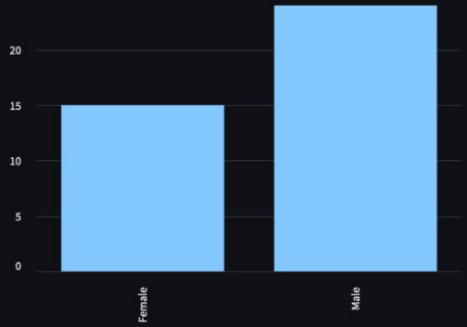

Diagnosis Distribution

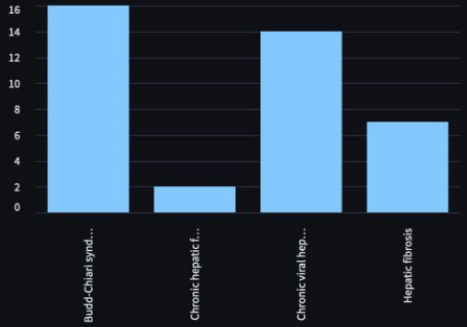

Age Distribution

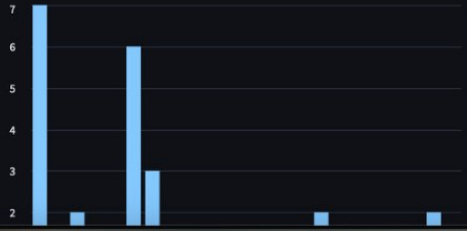

Python 3.10 (ollama-models-llm)

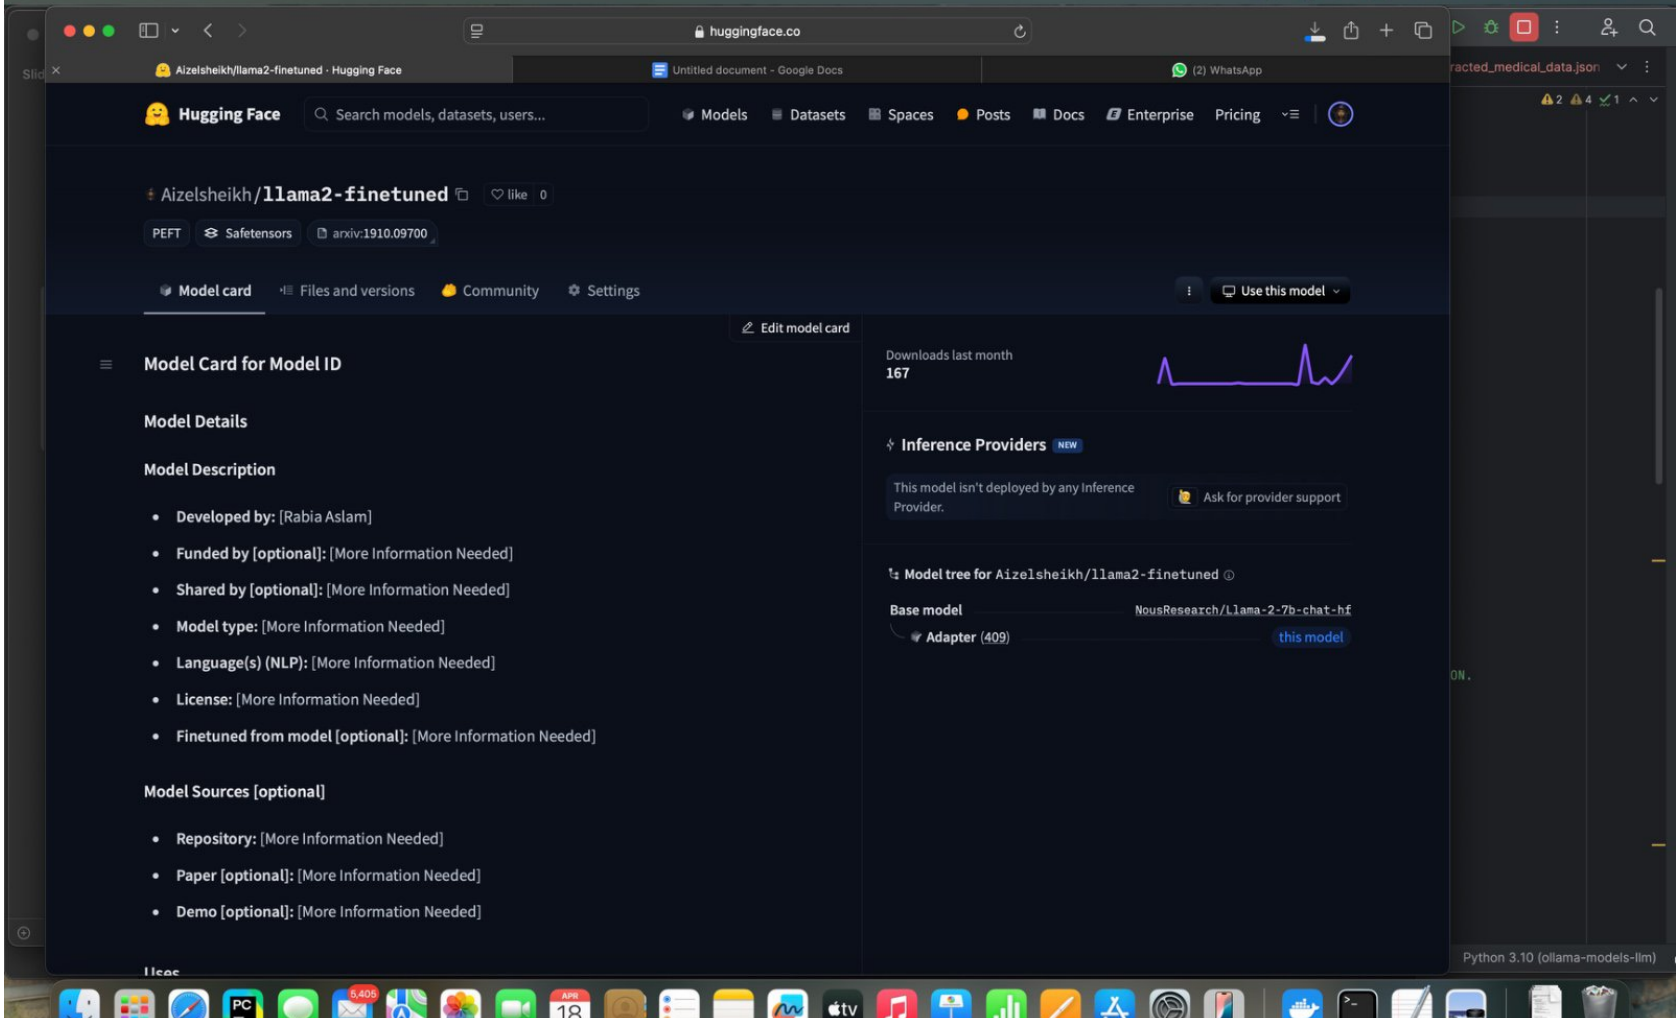
Safari File Edit View History Bookmarks Window Help
Fri 18 Apr 2:26 PM
huggingface.co
Aizelsheikh/llama2-finetuned · Hugging Face
Untitled document - Google Docs
(2) WhatsApp
Hugging Face
Search models, datasets, users...
Models Datasets Spaces Posts Docs Enterprise Pricing
Aizelsheikh/llama2-finetuned
like 0
PEFT
Safetensors
arxiv:1910.09700
Model card
Files and versions
Community
Settings
Use this model
Edit model card
Model Card for Model ID
Model Details
Model Description
Developed by: [Rabia Aslam]
Funded by [optional]: [More Information Needed]
Shared by [optional]: [More Information Needed]
Model type: [More Information Needed]
Language(s) (NLP): [More Information Needed]
License: [More Information Needed]
Finetuned from model [optional]: [More Information Needed]
Model Sources [optional]
Repository: [More Information Needed]
Paper [optional]: [More Information Needed]
Demo [optional]: [More Information Needed]
Downloads last month
167
Inference Providers NEW
This model isn't deployed by any Inference Provider.
Ask for provider support
Model tree for Aizelsheikh/llama2-finetuned
Base model
NousResearch/Llama-2-7b-chat-hf
Adapter (409)
this model
Python 3.10 (ollama-models-llm)